

CONTENT CLASSIFICATION FRAMEWORK

A Privacy-Preserving Alternative to
Identity-Based Age Verification

Design Specification

Version 1.1 • April 2026

DRAFT: For Discussion

Paul Moore

Information Security Consultant

<https://paul.reviews>

1. Executive Summary

Current approaches to protecting children from harmful online content rely on verifying the identity or age of users at the point of content consumption. This document presents evidence that such approaches are architecturally incapable of achieving their stated goal, and proposes an alternative framework that is simpler, more effective, more private, and deployable today.

The Content Classification Framework (CCF) inverts the data flow. Instead of requiring users to prove who they are, it requires websites to classify what they serve. The device, configured by a parent, device manufacturer, or operating system, makes local decisions about what content to display. No user data leaves the device. No identity credentials are required. No government infrastructure is needed. Adults are entirely unaffected.

This approach is modelled on a proven paradigm: television content ratings and the V-chip, which have operated successfully for decades without requiring viewers to identify themselves. The CCF adapts this model for the internet, with granularity that television could never achieve: per-page, per-element, and per-segment classification.

2. Problem Statement

2.1 The Fundamental Impossibility

Anonymous age verification for content gating is a logical impossibility. This has been proven from first principles through a security audit of the EU Digital Identity Wallet age verification system and its 81 associated repositories.

The proof is straightforward. Any system that attempts to verify age at the point of content consumption must satisfy two requirements simultaneously:

Requirement 1 (Privacy): The verifier must not be able to track, identify, or correlate users across sessions. This is the GDPR obligation, the political promise, and the public expectation.

Requirement 2 (Verification): The verifier must confirm that the person at the screen is the person whose age was verified. This is what age verification actually requires.

These requirements are mutually exclusive. Any data that satisfies Requirement 2 violates Requirement 1. Any system that satisfies Requirement 1 enables a relay attack where an adult's credential is used on a child's device. The privacy model that prevents tracking is the same property that prevents the verifier from detecting the relay. This is not an implementation gap. It is a mathematical constraint that no protocol can overcome.

2.2 The Escalation Ratchet

When an identity-based age verification system fails to protect children (which it inevitably will, for the reasons stated above), the institutional response follows a predictable escalation:

Step 1: Anonymous age token. Fails because it cannot verify who is at the screen. Response: "We need more data."

Step 2: Full identity disclosure. Name, date of birth, passport number. Fails because a relay still transmits the full identity. Response: "We need real-time verification."

Step 3: Biometric check at point of consumption. Live camera, face match. Requires every user to submit biometrics to every age-restricted site. Fails because deepfakes, photos, and an adult verifying before handing to a child all defeat it. Response: “We need continuous verification.”

Step 4: Continuous biometric monitoring. Camera stays on, periodic re-checks, behavioural biometrics. Total surveillance of every user session. Still fails because it cannot determine age from appearance or behaviour.

Each step fails for the same structural reason. Each failure justifies the next escalation. Each escalation trades more privacy for the same ineffective outcome. The ratchet only turns in one direction.

2.3 The Wrong Layer

Identity-based age verification operates at the wrong architectural layer. It attempts to build a firewall around every website rather than protecting the asset: the child’s device. This is equivalent to requiring every shop in the world to verify every customer’s age, rather than parents supervising what their children purchase.

The internet is global, borderless, and largely outside any single jurisdiction’s enforcement capability. Mandating age verification at the platform layer affects only platforms within regulatory reach. Content hosted outside the EU, accessed via VPN, or served by non-compliant operators is entirely unaffected. The system gates a fraction of available content while imposing costs and privacy risks on every user.

3. Proposed Solution: Content Classification Framework

3.1 Core Principle

The Content Classification Framework inverts the current model:

Current model: The user tells the site something about themselves (age, identity, credential). The site decides whether to grant access.

Proposed model: The site tells the device something about its content (category, rating, risk level). The device decides whether to display it.

The information flow is reversed. User data never leaves the device. The site classifies its own content. The device, configured by a parent, device manufacturer, or operating system, enforces the appropriate rules locally. The site has no knowledge of whether its content was displayed, filtered, or blocked.

3.2 Precedent

This model has operated successfully in broadcast media for decades. Television content ratings (TV-Y, TV-PG, TV-MA in the US; U, PG, 12, 15, 18 in the UK) are embedded in the broadcast signal. The V-chip reads the rating and enforces parental configuration. Broadcasters have no knowledge of which households block their content. The decision is entirely local.

The CCF adapts this proven model for the internet, with three significant enhancements: per-page and per-element granularity (rather than per-channel), per-segment granularity for streaming content, and crowd-sourced verification of classification accuracy.

3.3 Sites Already Classify Their Content

Under the current regulatory model (both the UK Online Safety Act and the EU DSA), every website must already determine whether its content requires age verification. That determination is itself a content classification. The site has already assessed its own content against regulatory thresholds and decided: “this is adult content” or “this is not adult content.” The hard work is already done.

Having made that classification, the current model then requires the site to build or integrate an entire age verification system: credential checking, Digital Credentials API integration, trusted list validation, EUDI wallet support, third-party identity provider contracts, GDPR data processing agreements, breach response plans, and ongoing compliance monitoring. All of that infrastructure sits on top of a classification decision the site has already made.

The Content Classification Framework asks the site to express the classification it has already made as a response header. That is the entire obligation. The classification work is done. The only difference is the output format: a header instead of an age gate.

Furthermore, major platforms already classify content at a granular level for advertising purposes. Brand-safety classification systems used to sell targeted advertising categorise pages by content type, risk level, and audience suitability. The framework asks platforms to share with parents and devices the same content classification they already share with advertisers. This is not a new burden. It is the exposure of existing data in a standardised format.

3.4 Zero Impact on Adult Traffic

Identity-based age verification systems impose friction on every user. Each visitor must present a credential, scan a passport, or interact with a third-party identity provider before accessing content. This friction measurably reduces traffic. Users who are unwilling to identify themselves leave. Users who encounter errors in the verification flow leave. Users concerned about privacy leave. Each lost visitor represents lost advertising revenue, reduced engagement metrics, and diminished platform value.

The Content Classification Framework imposes zero friction on any user. Adults do not see it, interact with it, or know it exists. Their browsing experience is entirely unchanged. Traffic volumes, advertising delivery, engagement metrics, and conversion rates are unaffected. The only change is a response header that the adult’s browser silently ignores.

For platforms, this means content classification can be adopted without any measurable impact on revenue or user metrics. There is no business case against adoption because there is no business cost. The framework protects children without penalising the platforms that serve them or the adults who use them.

3.5 Design Philosophy

Privacy by architecture, not by policy. No user data flows to any external party. Privacy is not a feature that can be disabled by policy change. The data simply does not exist because it was never collected.

Proportionality. Adults are entirely unaffected. No credentials, no scanning, no extensions, no gates. The unrestricted experience is the default.

Shared responsibility. The framework does not place the burden solely on parents. Device manufacturers can ship devices with sensible defaults. Operating systems can integrate content classification enforcement natively. Parents can customise settings for their children. Schools and institutions can configure devices under their management. Each party contributes what it is best positioned to contribute.

Honesty about limitations. This framework does not solve the problem of children with completely disengaged parents using unmanaged devices. Nothing does. That is a child welfare problem, not a technology problem. The framework protects the children who can be protected without harming everyone else in the attempt.

No surveillance infrastructure. The framework creates nothing that can be repurposed for surveillance, censorship, or identity tracking. There is no centralised database, no credential infrastructure, and no escalation path.

4. Technical Specification

4.1 Content Classification Mechanisms

Sites classify their content through four complementary mechanisms, each operating at a different granularity level. These mechanisms are layered: the most specific classification takes precedence, with each layer providing a fallback for the one above it.

4.1.1 Domain-Level Classification: `.well-known/content-rating.json`

A static JSON file at a well-known URI provides the default content classification for an entire domain. This is analogous to `robots.txt`: a lightweight, universally understood mechanism that requires minimal implementation effort.

```
GET /.well-known/content-rating.json

{
  "version": "1.0",
  "default_rating": "safe",
  "categories": {
    "/news": "safe",
    "/opinion": "mature/political",
    "/gallery/nsfw": "adult/sexual",
    "/forums/18+": "adult/general"
  },
  "contact": "trust-safety@example.com",
  "last_updated": "2026-04-01"
}
```

The categories object allows path-prefix matching, enabling a single file to classify different sections of a site. If no path matches, the `default_rating` applies. This accommodates mixed-content sites such as Reddit, YouTube, and news organisations where the domain is safe but specific paths may not be.

4.1.2 Page-Level Classification: HTTP Response Header

An HTTP response header provides per-page classification that overrides the domain-level default. This is the primary mechanism for dynamic content and user-generated content platforms.

```
HTTP/1.1 200 OK
Content-Type: text/html
Content-Rating: adult/sexual
Content-Rating-Source: platform-classified
```

The Content-Rating header travels with each individual response. The server, which already knows what content it is serving at that specific URL, classifies it per response. This enables platforms with user-generated content to classify dynamically based on their existing internal content analysis systems: the same systems they already use for advertising brand-safety classification.

4.1.3 Element-Level Classification: Inline Content Attributes

Individual HTML elements can carry content classification attributes, enabling sub-page granularity. The site marks specific elements with a data-content-rating attribute:

```
<article data-content-rating="safe">
  <p>Educational content about conflict resolution...</p>
  <div data-content-rating="mature/violence">
    
  </div>
  <p>The peace process continued with...</p>
</article>
```

The browser or parental control system reads these attributes and applies local enforcement: blur, hide, collapse, or replace with a placeholder. The surrounding content remains visible. The page does not break. A child reads the article and sees a blurred image where the graphic photo was. Everything else renders normally.

This approach is self-contained. The classification travels with the content itself, not in a separate file that might fall out of sync. When a developer adds a new content section, they add the rating attribute at the same time, in the same markup. The data-content-rating attribute is valid HTML5, uses the standard data-* convention, and has zero effect on browsers or devices that do not process it. Adults see no change whatsoever.

Element-level classification composes naturally with the page-level header. The Content-Rating header provides the page default. Individual data-content-rating attributes provide overrides within the page. A page rated safe might contain one element rated mature/violence. The device applies the page default and then checks individual elements for overrides. Two layers, working together, enabling a news article to remain accessible while a single graphic image within it is filtered.

4.1.4 Segment-Level Classification: Streaming Metadata

For streaming content delivered via HLS or DASH, content ratings can be embedded at the segment level in the manifest file. This enables unprecedented granularity. A documentary

that is educational for 95% of its runtime but contains 30 seconds of graphic imagery can be classified per-segment.

```
#EXTINF:4.000,
#EXT-X-CONTENT-RATING:safe
segment_001.ts

#EXTINF:4.000,
#EXT-X-CONTENT-RATING:violence/graphic
segment_047.ts

#EXTINF:4.000,
#EXT-X-CONTENT-RATING:safe
segment_048.ts
```

The device's parental control system reads the manifest, identifies restricted segments, and applies the configured response: blur the video, mute the audio, skip the segment, or display a placeholder. The stream continues. The child watches the documentary and is shielded from the graphic segment. The streaming service has no knowledge that the device modified the playback.

This eliminates the binary choice that plagues current content moderation: block entire content or allow all of it. A child can watch an age-appropriate film and have a single scene filtered. A teenager can watch a news broadcast with graphic conflict footage blurred. Content remains available; the experience is adjusted.

4.2 Content Rating Taxonomy

The framework uses a hierarchical taxonomy that maps to established content rating systems (PEGI, ESRB, BBFC) while remaining jurisdiction-neutral. Sites classify their content; the device maps it to the configured rules based on local standards.

Category	Sub-category	Description
safe	(none)	Suitable for all audiences. No restrictions expected.
mature	language	Profanity, strong language
mature	violence	Depictions of violence, conflict imagery
mature	political	Politically sensitive or partisan content
mature	health	Medical, sexual health, or drug-related content
adult	sexual	Sexually explicit content
adult	gambling	Gambling, betting, or casino content
adult	substances	Content promoting alcohol, tobacco, or drug use
adult	general	General adult content not otherwise categorised

The taxonomy is intentionally extensible. New categories can be added without breaking existing implementations. Devices that do not recognise a category can fall back to the top-level classification (safe, mature, adult).

The configuration offers granular control. A parent may permit mature/health content for their 15-year-old (allowing access to sexual health resources) while blocking mature/violence. A device manufacturer may ship with sensible defaults that block adult categories out of the box. An operating system may integrate the taxonomy into its existing parental control interface. A school may configure all devices under its management to block mature and adult categories during school hours. The framework enables the decision; it does not make the decision.

4.3 Device-Side Enforcement

Content classifications are consumed and enforced entirely on the device. The enforcement mechanism depends on the integration level:

Device manufacturer defaults: Manufacturers can ship devices with content classification enforcement enabled by default, particularly for devices marketed to families or children. A tablet sold as a children's device can enforce age-appropriate content out of the box, with no configuration required from the parent. The parent retains the ability to adjust settings, but the default is safe.

Operating system integration: Android, iOS, Windows, and ChromeOS already have parental control frameworks (Family Link, Screen Time, Microsoft Family Safety). Content classification enforcement integrates natively with these systems. The OS reads content ratings and enforces them based on the family management configuration. This is architecturally identical to how operating systems already enforce app age ratings from the Play Store and App Store.

Browser-native integration: Chromium, WebKit, and Gecko read the Content-Rating header, data-content-rating attributes, and the .well-known file natively. The browser checks classifications against the OS-level parental control configuration. If the content exceeds the configured threshold, the browser blocks the page, blurs elements, displays a placeholder, or applies content filtering as appropriate to the classification level.

OS-level media framework: For streaming content, the media playback framework (ExoPlayer, AVFoundation) reads segment-level ratings from the manifest. Restricted segments are blurred, muted, skipped, or replaced with a placeholder, based on the device configuration.

Parental control applications: Third-party parental control software can read content classifications via system APIs exposed by the browser. This allows existing parental control ecosystems to use the framework without replacement.

Browser extension (proof of concept): A browser extension can demonstrate the framework's feasibility today, without browser-level or OS-level integration. A proof-of-concept extension already exists and is functional. This serves as a deployment bridge while native integration is developed.

4.4 Handling Unclassified Content

When no classification exists for a given page (no header, no .well-known, no community classification list entry), the device's response should be configurable:

Allow by default: Suitable for general-purpose devices. Unclassified content is treated as unrestricted. This is the least disruptive default and appropriate for most adult users.

Warn and allow: The device flags unclassified content with a visual indicator but does not block it. The parent or user is informed that no classification was available.

Block by default: Suitable for devices configured for young children. Unclassified content is blocked until a parent explicitly approves it. This is the most restrictive default.

Device manufacturers choose the appropriate default for their target market. A general-purpose smartphone defaults to allow. A children's tablet defaults to block. A school Chromebook defaults to warn. The parent can override any default. This turns the absence of a classification into a configurable policy rather than a gap in protection.

4.5 Third-Party Content and Iframes

A page may be classified as safe but embed an iframe from an unsafe source. Each iframe's content carries its own classification from its own origin. The browser already enforces same-origin policy across iframe boundaries. Extending this to content classification is architecturally natural.

A safe page with an unclassified or adult-rated iframe has the iframe blocked or blurred while the parent page renders normally. This prevents the common bypass where a safe domain loads restricted content through a third-party embed. The classification of each resource is evaluated independently based on its own origin's classification.

4.6 On-Device ML Classification as Fallback

For sites that do not self-classify and are not covered by any community classification list, the device can run lightweight on-device machine learning models to generate a local classification estimate. No data leaves the device. The model runs locally. The classification is imperfect but provides a safety net for content that would otherwise be entirely unclassified.

This is already technically feasible. Browsers run ML models on-device for features such as phishing detection, image alt-text generation, and safe browsing warnings. Extending this capability to content classification is a natural evolution. The on-device model serves as the final fallback in the classification hierarchy: site self-classification takes precedence, followed by community lists, with on-device ML as the last resort.

The output of the on-device model is treated as advisory, not authoritative. It informs the device's filtering decision but can be overridden by the parent's configuration. If the model misclassifies a site, the parent can approve it manually. This avoids the false-positive problem that plagues automated content filters while still providing baseline protection for unclassified content.

4.7 No Data Leaves the Device

The privacy properties of this framework are not achieved through encryption, zero-knowledge proofs, or anonymisation techniques. They are achieved through architecture. No user data is collected, transmitted, or processed by any external party because the information flow is unidirectional: from site to device, never from device to site.

The site does not know whether the visitor is a child or an adult. It does not know whether parental controls are active. It does not know whether its content was displayed, filtered, or blocked. It does not know whether the device read the content classification at all.

This is a stronger privacy guarantee than any identity-based system can offer, regardless of the cryptographic techniques employed. The strongest form of data protection is never creating the data in the first place.

5. Crowd-Sourced Trust Intelligence

5.1 The Classification Accuracy Problem

Content classifications are self-reported by sites. A site that hosts adult content could classify itself as safe. The framework addresses this through multiple layers: crowd-sourced verification, community classification lists, and on-device ML fallback.

5.2 Community Classification Lists

The framework supports community-maintained classification lists, analogous to ad-blocker filter lists such as EasyList and EasyPrivacy. Volunteers and organisations curate lists of domain and path classifications. Devices subscribe to these lists. A site that does not self-classify still receives a classification from the community.

This means the framework provides value even before a single site participates. A parent installs the browser extension or configures their device today, subscribes to a classification list, and receives immediate protection. The framework does not depend on universal site adoption to be useful. This is a critical deployment advantage that identity-based systems cannot match, as they require platform integration before providing any value.

Multiple classification lists can coexist, maintained by different organisations with different perspectives. A parent chooses which lists to subscribe to, just as they choose which content filters to enable. This avoids centralised control over classification while enabling a rich ecosystem of community-driven content assessment.

5.3 Parent and Child Reporting

Parents and children who encounter misclassified content can report it through the parental control interface. These reports are anonymous: they identify the domain and the discrepancy, not the reporter. Reports aggregate into a trust intelligence layer that informs, but does not dictate, decisions.

When a child requests a site, the parental control system can indicate the community's assessment: "12 parents reported concerns about this site in the last week." The parent or device configuration retains the decision. The system provides context.

5.4 Abuse Prevention

Coordinated mass-reporting to suppress legitimate content (such as sex education resources, LGBT support, or political sites) is detectable through pattern analysis. Organic reports trickle in from diverse sources, geographies, and timeframes. Coordinated attacks cluster. AI-driven anomaly detection flags suspicious report patterns for review before they affect trust scores.

5.5 Domain Ownership Changes

Certificate Transparency logs, WHOIS records, and DNS changes are monitored automatically. When a domain changes ownership, its trust history resets. A previously

trusted domain that is sold and repurposed cannot trade on its prior reputation. This is an objective, automatable trigger that does not depend on user reports.

6. Deployment Options

6.1 Phase 1: Browser Extension and Community Lists (Available Now)

A browser extension provides immediate deployment capability without requiring browser vendor or OS vendor cooperation. Combined with community classification lists, this provides functional protection from day one, even without any site participation. A proof-of-concept extension already exists and is functional.

This phase validates the concept, builds a user base, and generates the adoption data needed to justify native integration.

6.2 Phase 2: Device Manufacturer and OS Integration

Device manufacturers and operating system vendors integrate content classification enforcement into their existing parental control frameworks. Android Family Link, iOS Screen Time, and Microsoft Family Safety already manage app-level age ratings. Extending these systems to read web content classifications is a natural evolution of existing capability, not a new product.

Manufacturers of devices marketed to children or families can enable content classification enforcement by default, providing out-of-the-box protection without requiring any parental action. This addresses the adoption gap where engaged parents already use parental controls and disengaged parents do not: the device itself provides the baseline protection.

6.3 Phase 3: Browser-Native Integration

The Content-Rating header is proposed as a web standard via the W3C or IETF. Chromium, WebKit, and Gecko implement native header and attribute parsing. The browser reads parental control configuration from the operating system's existing family management APIs.

This phase provides the trust property that a Sec- prefixed header offers: the header is processed by the browser engine, not by user-installable code that could be bypassed.

6.4 Phase 4: Streaming and Media Integration

Media frameworks (ExoPlayer, AVFoundation, Media Foundation) implement segment-level content rating parsing. Streaming services embed ratings in HLS/DASH manifests. The media player applies per-segment filtering based on the device configuration.

6.5 Phase 5: Ecosystem Maturity

Smart TVs, gaming consoles, IoT devices, and embedded browsers implement content classification parsing. The crowd-sourced trust layer and community lists mature. Search engines surface classification status as a trust signal, creating market incentives for site participation. The framework becomes a standard component of device-level content management across all platforms.

7. Adoption Incentives

7.1 Classification Status as a Trust Signal

Search engines can surface a site's classification status in search results. A site that self-classifies its content could display a trust badge indicating participation in the framework. Parents searching for child-appropriate content can filter by classification status, preferring sites that have opted in.

This creates a positive market incentive to participate. Sites that classify their content are rewarded with increased visibility and trust. Sites that do not are not penalised, but they miss the opportunity. The incentive is additive rather than punitive, encouraging voluntary adoption without requiring regulatory mandates.

7.2 Platform Business Case

The business case for platform adoption is straightforward: zero cost, zero risk, positive positioning. Adding a Content-Rating header or publishing a .well-known/content-rating.json file requires minimal engineering effort. No user data is processed, so there is no GDPR liability. No users are blocked, so there is no traffic loss. No credentials are checked, so there is no integration cost.

Platforms that adopt the framework can publicly state their commitment to child safety without imposing any friction on their users. This is a marketing advantage with no corresponding downside. The framework is the rare regulatory alignment where compliance is genuinely costless.

8. Regulatory Alignment

8.1 DSA Compliance

The Digital Services Act (Article 28) requires platforms to implement appropriate and proportionate measures to ensure a high level of privacy, safety, and security of minors on their services. Content classification is both appropriate and proportionate:

Appropriate because it provides platforms with a mechanism to support age-appropriate design without collecting user identity data. Platforms already classify their content internally for advertising brand-safety purposes. The framework asks them to expose the same classification in a machine-readable format.

Proportionate because it imposes minimal burden on platforms (a response header or a .well-known file), no burden on adult users (the default experience is unchanged), and no privacy cost to any user (no data leaves the device). It is the least invasive measure capable of achieving meaningful child protection.

8.2 Sites Already Do the Hard Work

Under both the UK Online Safety Act and the EU DSA, websites must already determine whether their content requires age verification. This determination is itself a content classification. The site has already assessed its content against regulatory thresholds.

The current model then requires the site to build or integrate an entire age verification infrastructure on top of that classification. The Content Classification Framework asks the site to express that same classification as a response header. The classification work is already done. The only change is the output format.

For major platforms, the classification is even more granular than regulatory thresholds require. Brand-safety systems used for advertising already categorise individual pages by content type, risk level, and audience suitability. The framework asks platforms to share with devices the same classification they already share with advertisers. This is not a new burden. It is the standardised exposure of existing data.

8.3 GDPR Compatibility

The framework processes no personal data whatsoever. The content classification is information about the site, not about the user. The device-side enforcement processes no data that leaves the device. There is no data controller, no data processor, no data subject, and no processing activity within the meaning of the GDPR. The framework operates entirely outside the scope of data protection regulation because it creates no data protection risk.

8.4 Enforcement

A site that misclassifies its content, declaring itself as safe when it hosts adult material, is making a false declaration. This is enforceable through existing consumer protection and advertising standards law in most jurisdictions. No new legislation is required.

The crowd-sourced trust layer and community classification lists provide the evidence base for enforcement. A site with hundreds of parent reports disputing its self-classification is a clear signal for regulatory investigation, without requiring any centralised monitoring infrastructure.

9. Anti-Censorship Safeguards

The architecture that enables parental content control is structurally similar to the architecture that enables censorship. The difference is governance, not technology. This section codifies the safeguards that prevent the framework from being repurposed as a censorship tool. These safeguards are normative requirements of the specification, not optional recommendations. Any implementation that violates them is non-compliant.

9.1 Lists Are Optional, Never Mandated

No government or authority can require device manufacturers or operating systems to subscribe to a specific classification list. Parents choose their lists. Manufacturers choose their default lists. Governments can publish lists, but subscription is always voluntary. The specification explicitly prohibits mandated list subscription. The moment a government can compel subscription to its list, the framework becomes a censorship infrastructure. This principle is non-negotiable.

9.2 Lists Are Transparent and Auditable

Every classification list is published in full. Every entry is visible to anyone. There are no secret classifications. A government list that classifies a political opposition site as adult/sexual is visible to journalists, civil liberties organisations, and the public. The specification explicitly prohibits secret or partially redacted lists. Transparency is the primary defence against abuse.

9.3 Self-Classification Takes Precedence

If a site self-classifies as safe and a third-party list classifies it differently, the device surfaces the conflict to the user rather than silently applying either classification. The parent sees both

classifications and makes their own decision. Conflicts are surfaced, never resolved silently. This prevents a list from overriding a site's honest self-classification without the user's knowledge.

9.4 No List Can Block Access Entirely

Lists classify content. They do not block domains. The device's response to a classification is always configurable by the parent. Even on the most restrictive settings, a parent can explicitly approve any site. There is no technical mechanism for a list to make content permanently inaccessible. A government that wants to block a site must use DNS or network-level blocking, which is visible and subject to legal challenge. They cannot use the CCF to achieve invisible blocking.

9.5 Adults Are Never Affected

This is the architectural guarantee against censorship at scale. The framework only operates on devices with parental controls or child profiles active. An adult's device ignores all classification lists entirely. There is no mechanism by which a government could use the CCF to restrict what adults see. If a government wanted to censor content for the general population, they would need an entirely different infrastructure. The CCF is architecturally incapable of serving that purpose.

9.6 Multiple Competing Lists Prevent Monopoly

The specification explicitly supports multiple concurrent list subscriptions. A parent can subscribe to their government's list, an NGO's list, a religious organisation's list, and a civil liberties organisation's list simultaneously. Conflicts between lists are surfaced to the parent. No single entity controls the classification landscape. This mirrors the ad-blocker ecosystem where EasyList, EasyPrivacy, and regional lists coexist, and the user chooses which to enable.

9.7 Community Reporting Applies to Lists

If a classification list consistently misclassifies content (for example, a government list that systematically classifies opposition media as adult content), users can report the list itself. List reputation becomes part of the trust intelligence layer. A list with a pattern of politically motivated classifications accumulates negative trust signals. Parents are warned. Journalists investigate. Lists are held accountable by the same crowd-sourced mechanisms that hold sites accountable.

The honest assessment is that the CCF cannot be made completely censorship-proof, because any content classification system can be abused by a sufficiently determined government. But it can be made censorship-resistant through transparency, optionality, competing lists, conflict surfacing, and the fundamental architectural constraint that it only operates on children's devices. A government that wants to censor adults cannot use this framework to do it. That alone makes it categorically different from identity-based age verification, where the infrastructure operates on every user and can be repurposed with a policy change.

10. Comparison with Identity-Based Age Verification

Property	Identity-Based AV	Content Classification
User data collected	Passport, biometrics, credentials	None
Adult impact	Must present ID/credential	None; unaware of framework
Relay attack	Architecturally unsolvable	Not applicable; no credential
Platform cost	Significant (integration, liability)	Minimal (response header)
Traffic impact	Measurable loss (verification friction)	Zero (adults unaffected)
GDPR exposure	Full (sensitive data processing)	None (no personal data)
Surveillance risk	Infrastructure adaptable	No infrastructure to repurpose
Content granularity	Binary (access/deny per site)	Per-page, per-element, per-segment
VPN bypass	Defeated entirely	Unaffected (device-side)
Deployment complexity	81 repositories, 7-country pilot	1 header, 1 .well-known file
Day-one value	Requires platform integration	Community lists work immediately
Breach risk	Catastrophic (passport data)	None (no data held)
Escalation path	Inevitable (more data needed)	None (no data to expand)

11. Scope and Limitations

This document intentionally presents an honest assessment of what the Content Classification Framework can and cannot achieve.

It can: Provide parents, device manufacturers, and operating systems with effective, granular, privacy-preserving mechanisms to manage children’s online experience. Give platforms a low-cost, zero-liability mechanism to support age-appropriate design with zero impact on adult traffic. Eliminate the need for identity-based age verification and the surveillance infrastructure it requires. Provide per-page, per-element, and per-segment content adjustment rather than binary access/deny decisions. Enable manufacturers to ship devices with safe defaults that protect children out of the box. Provide value from day one through community classification lists, even without site participation.

It cannot: Protect children on unmanaged devices with no parental controls and no manufacturer defaults. Force sites to classify their content honestly without enforcement. Prevent a technically sophisticated teenager from circumventing device-level controls on a device they fully control.

These limitations are shared by every child safety system. The difference is that the CCF does not pretend to overcome them and therefore does not generate pressure for escalation when they manifest. A system that promises to protect all children and fails will be expanded. A system that provides effective tools to parents, manufacturers, and platforms, and succeeds within its stated scope, will be stable.

12. Next Steps

The following actions are recommended to advance the Content Classification Framework from concept to deployment:

- 1. Standards proposal:** Draft a W3C Community Group report or IETF Internet-Draft defining the Content-Rating response header, the data-content-rating HTML attribute, the .well-known/content-rating.json format, and the HLS/DASH segment-level rating extension.
- 2. Community classification lists:** Establish the initial community classification list infrastructure. Engage existing content safety organisations to seed the first lists. Define the list format and subscription mechanism.
- 3. Device manufacturer engagement:** Present the proposal to major device manufacturers (Samsung, Google, Apple, Microsoft) as a native integration for their existing parental control frameworks. Propose default-on enforcement for devices marketed to families and children.
- 4. Browser vendor engagement:** Present the proposal to Chromium, WebKit, and Gecko teams. Request integration with existing OS-level parental control APIs (Family Link, Screen Time, Microsoft Family Safety).
- 5. Platform pilot:** Engage 2 to 3 major platforms (ideally one social media, one streaming, one news) to implement content classifications. Measure classification accuracy and content management effectiveness.
- 6. Crowd-sourced trust MVP:** Build the minimum viable reporting and trust intelligence layer. Deploy alongside the browser extension for early validation.
- 7. Regulatory engagement:** Present the framework to the European Commission as a compliant alternative under DSA Article 28. Position content classification as a proportionate measure that achieves age-appropriate design without identity verification.
- 8. Streaming integration proposal:** Engage HLS (Apple) and DASH (MPEG) standards bodies to propose segment-level content rating extensions to existing manifest specifications.

Content Classification Framework v1.1

April 2026 | Draft for Discussion

Paul Moore | <https://paul.reviews>